

NoSQL Databases

PL05 – Aggregation and Indexing in
MongoDB

Teacher: Cristiana Neto

Email: cristiana.neto@algoritmi.uminho.pt

Office hours:

Friday 10h–11h



Summary

1

Aggregation in MongoDB

2

Indexing

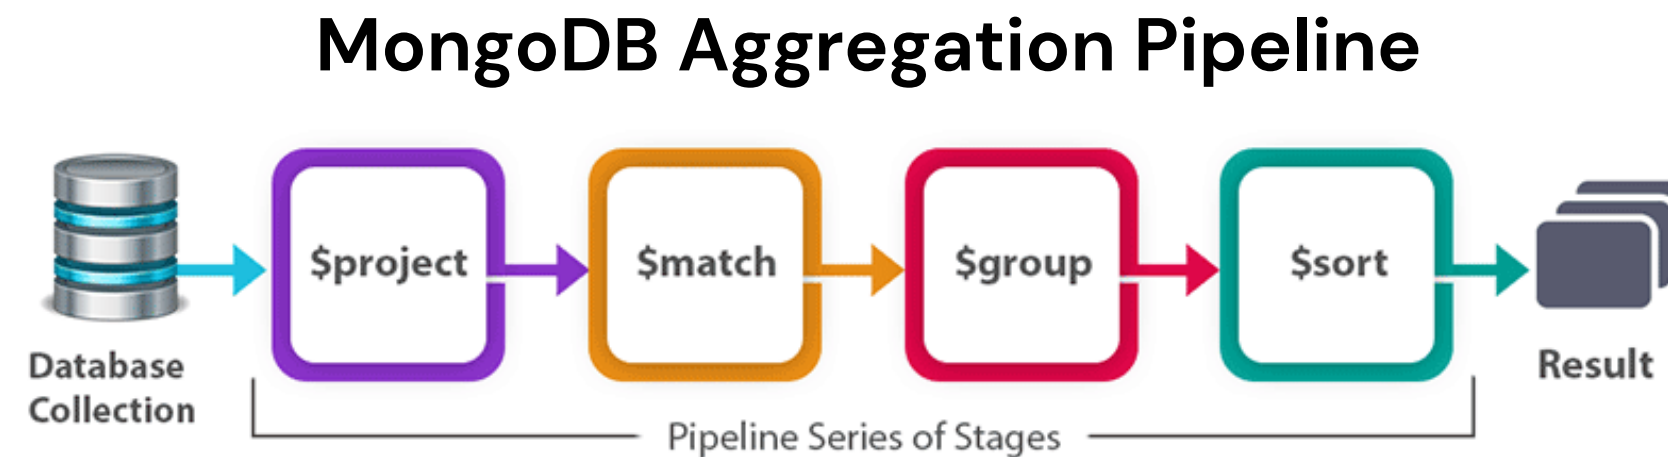
3

FE04 – Worksheet 4

Aggregation

➔ Pipeline

Aggregation operations in MongoDB are operations that group data together and return computed results, such as sums and counts.



`$project` – Select the document fields we want to work with

`$match` – It filters the documents we need to work with, which suit our needs

`$group` – does the aggregation work

`$sort` – sorts the resulting documents the way we want (descending or ascending)

Aggregation

➔ Aggregation Syntax

```
db.collectionName.aggregate(pipeline, options)
```

- *collectionName* – Collection Name
- *pipeline* – Array that contains the steps
- *options* – Optional parameters for aggregation

```
pipeline = [  
    { $match : { ... } },  
    { $group : { ... } },  
    { $sort : { ... } }  
]
```

Aggregation

➔ Example #1

Create db:

```
> use p105
```

Create Universities Collection and Add Documents:

```
> db.universities.insertMany([ { country : 'Spain', city : 'Salamanca', name :  
'USAL', location : { type : 'Point', coordinates : [ -5.6722512,17,  
40.9607792 ] }, students : [ { year : 2014, number : 24774 }, { year : 2015,  
number : 23166 }, { year : 2016, number : 21913 }, { year : 2017, number :  
21715 } ] }, { country : 'Spain', city : 'Salamanca', name : 'UPSA', location :  
{ type : 'Point', coordinates : [ -5.6691191,17, 40.9631732 ] }, students :  
[ { year : 2014, number : 4788 }, { year : 2015, number : 4821 }, { year :  
2016, number : 6550 }, { year : 2017, number : 6125 } ] } ])
```

Create courses collection and add documents:

```
> db.courses.insertMany([ { university : 'USAL', name : 'Computer Science', level  
: 'Excellent' }, { university : 'USAL', name : 'Electronics', level :  
'Intermediate' }, { university : 'USAL', name : 'Communication', level :  
'Excellent' } ])
```

Aggregation

➔ Example #1

3. Save the results of an aggregation to a new collection:

```
db.universities.aggregate([
  { $group: { _id: '$name', totaldocs: { $sum: 1 } } },
  { $out: 'aggResults' }
])
```

4. Obtain a document for each element of the USAL university student array:

```
db.universities.aggregate([
  { $match : { name : 'USAL' } },
  { $unwind : '$students' } ])
```

Aggregation

➔ Example #1

5. Sort the documents obtained in the previous point by the number of students in descending order, projecting only the year and the number of students:

```
db.universities.aggregate([
  { $match : { name : 'USAL' } },
  { $unwind : '$students' },
  { $project : { _id : 0, 'students.year' : 1, 'students.number' : 1 } },
  { $sort : { 'students.number' : -1 } } ])
```

Aggregation

➔ Example #1

6. Limit Search:

```
db.universities.aggregate([
  { $match : { name : 'USAL' } },
  { $unwind : '$students' },
  { $project : { _id : 0, 'students.year' : 1, 'students.number' : 1 } },
  { $sort : { 'students.number' : -1 } },
  { $limit : 1 }
])
```


Aggregation

➔ Example #1

7. Check the amount of documents obtained in the output of the previous steps of the pipeline:

```
db.universities.aggregate([  
    { $unwind : '$students' },  
    { $count : 'total_documents' } ])
```

Aggregation

➔ Example #1

8. Merge fields from two collections:

```
db.universities.aggregate([
  { $match : { name : 'USAL' } },
  { $project : { _id : 0, name : 1 } },
  { $lookup :
    { from : 'courses',
      localField : 'name',
      foreignField : 'university',
      as : 'courses' }
  } ])
```

PS: If we want this query to run quickly, we will need to index the `name` field in the `universities` collection and the `university` field in the `courses` collection.

Aggregation

➔ Example #1

9. Total number of students who have belonged to each of the universities in descending order:

```
db.universities.aggregate([
  { $unwind : '$students' },
  { $group : { _id : '$name', totalalumni : { $sum : '$students.number' } } },
  { $sort : { totalalumni : -1 } }
])
```

Aggregation

➔ Example #2

Create db:

```
> use pizza
```

Create order collection and add documents:

```
> db.orders.insertMany( [  
  { _id: 0, name: "Pepperoni", size: "small", price: 19, quantity: 10, date: ISODate("2021-03-13T08:14:30Z") },  
  { _id: 1, name: "Pepperoni", size: "medium", price: 20, quantity: 20, date: ISODate("2021-03-13T09:13:24Z") },  
  { _id: 2, name: "Pepperoni", size: "large", price: 21, quantity: 30, date: ISODate("2021-03-17T09:22:12Z") },  
  { _id: 3, name: "Cheese", size: "small", price: 12, quantity: 15, date: ISODate("2021-03-13T11:21:39.736Z") },  
  { _id: 4, name: "Cheese", size: "medium", price: 13, quantity: 50, date: ISODate("2022-01-12T21:23:13.331Z") },  
  { _id: 5, name: "Cheese", size: "large", price: 14, quantity: 10, date: ISODate("2022-01-12T05:08:13Z") },  
  { _id: 6, name: "Vegan", size: "small", price: 17, quantity: 10, date: ISODate("2021-01-13T05:08:13Z") },  
  { _id: 7, name: "Vegan", size: "medium", price: 18, quantity: 10, date: ISODate("2021-01-13T05:10:13Z") }  
)
```

Aggregation

➔ Example #2

1. Return the total order quantity of medium pizzas grouped by pizza name:

```
> db.orders.aggregate( [  
  // Stage 1: Filter your order documents by pizza size  
  {  
    $match: { size: "medium" }  
  },  
  // Stage 2: Group the remaining documents by the name of the pie and calculate  
  the total quantity  
  {  
    $group: { _id: "$name", totalQuantity: { $sum: "$quantity" } }  
  }  
] )
```

Result:

```
[  
  { _id: 'Cheese', totalQuantity: 50 },  
  { _id: 'Vegan', totalQuantity: 10 },  
  { _id: 'Pepperoni', totalQuantity: 20 }  
]
```

Aggregation

➔ Example #2

2. Return the total value and average order quantity between two dates:

```
> db.orders.aggregate( [  
  // Stage 1: Filter order documents by dates  
  {  
    $match: {"date": {$gte: new ISODate("2020-01-30"), $lt: new ISODate("2022-01-30")}}  
  },  
  // Stage 2: Group the remaining documents by date and calculate the results  
  {  
    $group: {_id: {$dateToString: {format: "%Y-%m-%d", date: "$date" } },  
      totalOrderValue: {$sum: {$multiply: ["$price", "$quantity"] } },  
      averageOrderQuantity: {$avg: "$quantity" }},  
  },  
  // Stage 3: Sort documents by totalOrderValue in descending order  
  {  
    $sort: {totalOrderValue: -1}}  
  ] )
```

Aggregation in Compass

uni.universities

2 DOCUMENTS 1 INDEXES

Documents **Aggregations** Schema Explain Plan Indexes Validation

Pipeline  Your pipeline is currently empty. To get started add the [first stage](#).

Explain

Export

Run

[More Options ▶](#)

Untitled

 SAVE ▼

+ CREATE NEW

 EXPORT TO LANGUAGE

☒ PREVIEW

 STAGES

 TEXT



▼ 2 Documents in the collection 

Preview of documents

```
_id: ObjectId('6425af3de1a7114c3c60e6c2')
country: "Spain"
city: "Salamanca"
name: "USAL"
▶ location: Object
▶ students: Array
```

```
_id: ObjectId('6425af3de1a7114c3c60e6c3')
country: "Spain"
city: "Salamanca"
name: "UPSA"
▶ location: Object
▶ students: Array
```

+ Add Stage

[Learn more about aggregation pipeline stages](#) 

Indexing

Indexes support the efficient execution of queries in MongoDB. Without indexes, MongoDB must perform a *collection scan*, i.e. scan every document in a collection, to select those documents that match the query statement. If an appropriate index exists for a query, MongoDB can use the index to limit the number of documents it must inspect.

```
db.collection.createIndex( <key and index type specification>, <options> )
```

There are several types of indexes in MongoDB, including simple, compound, text, and geospatial indexes. Simple indexes index a single field, while compound indexes index multiple fields together. Text indexes are used for full-text search queries, while geospatial indexes are used for location-based queries.

FE04 – Aggregation in MongoDB



Universidade do Minho
Departamento de Informática

Curso: Mestrado em Informática / Mestrado em Bioinformática
U.C.: Bases de Dados NoSQL

Ficha de Exercícios 04 – PLO5	
Docente:	António Abelha / Cristiana Neto
Tema:	Agregação no MongoDB
Turma:	Prática Laboratorial
Ano Letivo:	2023-2024 – 2º Semestre
Duração da aula:	2 horas

FE04

NoSQL Databases

PL05 – Aggregation and Indexing in
MongoDB

Teacher: Cristiana Neto

Email: cristiana.neto@algoritmi.uminho.pt

Office hours:

Friday 10h–11h

